

Recurrent Neural Networks

Machine Learning for Behavioral Data

April 11, 2022

Today's Topic

Week	Lecture/Lab
1	Introduction
2	Data Exploration
3	Regression
4	Classification
5	Model Evaluation
6	Knowledge Tracing
7	Knowledge Tracing
8	Time Series Prediction

Supervised learning on time series:

- Probabilistic graphical models
- Neural networks: LSTM, GRU, etc.

Getting ready for today's lecture...

- **If not done yet:** clone the repository containing the Jupyter notebook and data for today's lecture into Google Colab
 - Join us on Kahoot for the Easter quiz!
-

Easter quiz about the past



www.kahoot.it

Enter the game pin!



Win a chocolate Easter bunny!



Today's Topic

Week	Lecture/Lab
1	Introduction
2	Data Exploration
3	Regression
4	Classification
5	Model Evaluation
6	Knowledge Tracing
7	Knowledge Tracing
8	Time Series Prediction

Supervised learning on time series:

- Probabilistic graphical models
- Neural networks: LSTM, GRU, etc.

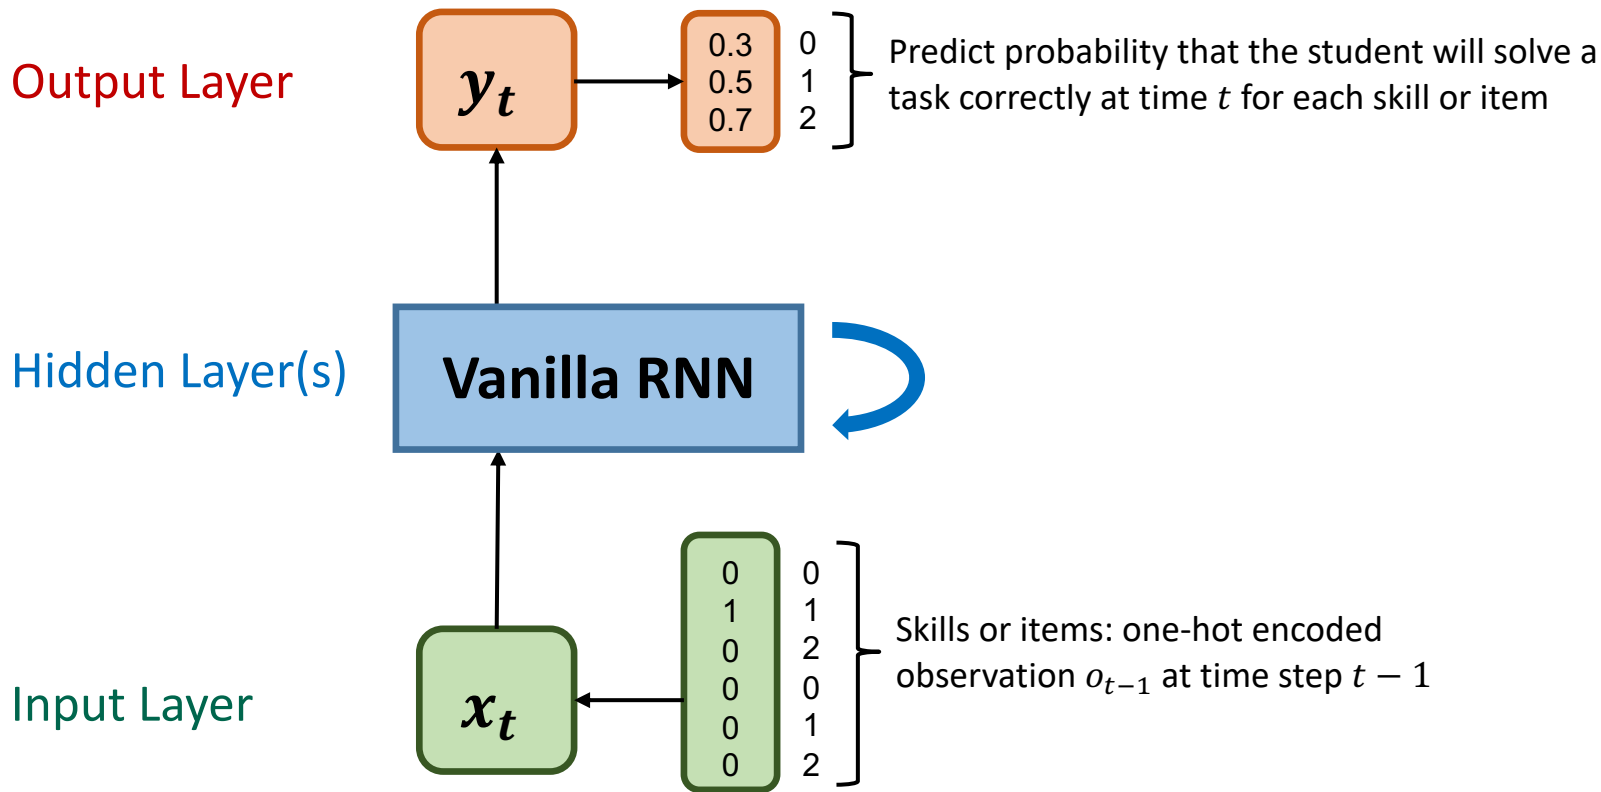
Today – Recurrent Neural Networks

- Parameters and hyperparameter tuning
 - Different architectures
 - Different tasks:
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression
-

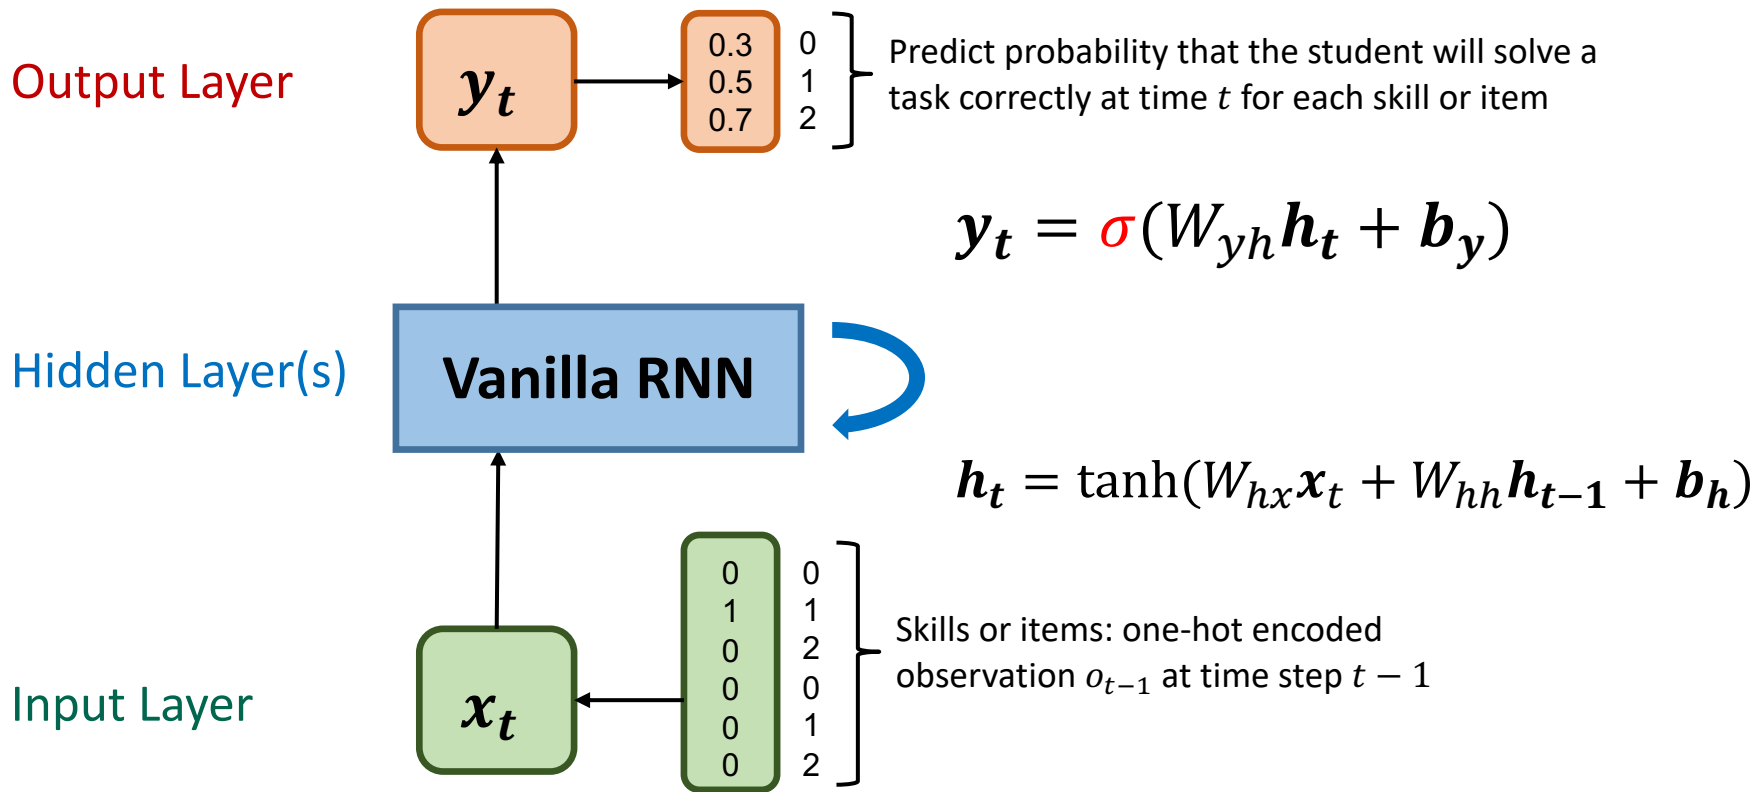
Today – Recurrent Neural Networks

- **Parameters and hyperparameter tuning**
 - Different architectures
 - Different tasks:
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression
-


Deep Knowledge Tracing - revisited




Deep Knowledge Tracing - revisited



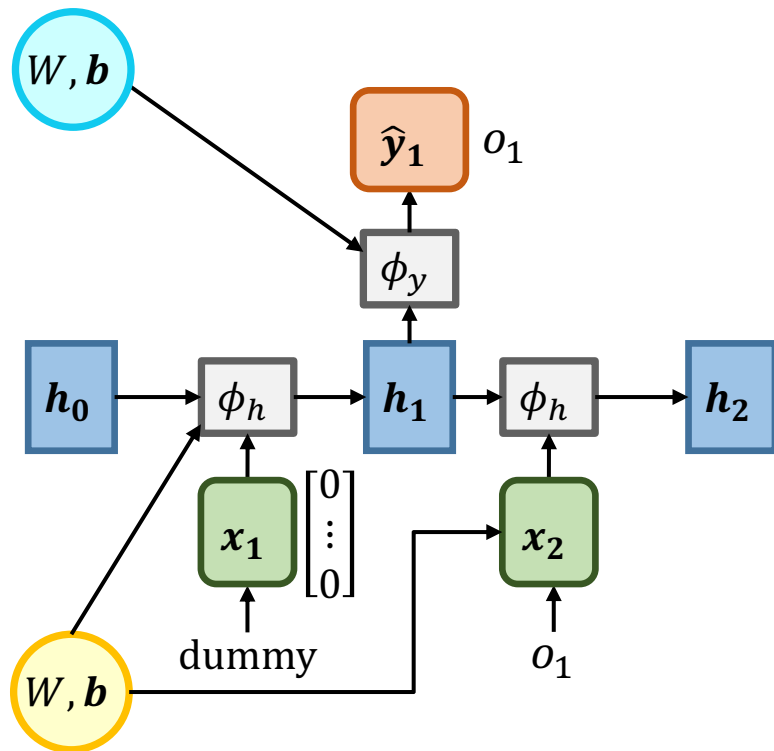
Computational Graph - revisited




Computational Graph - revisited



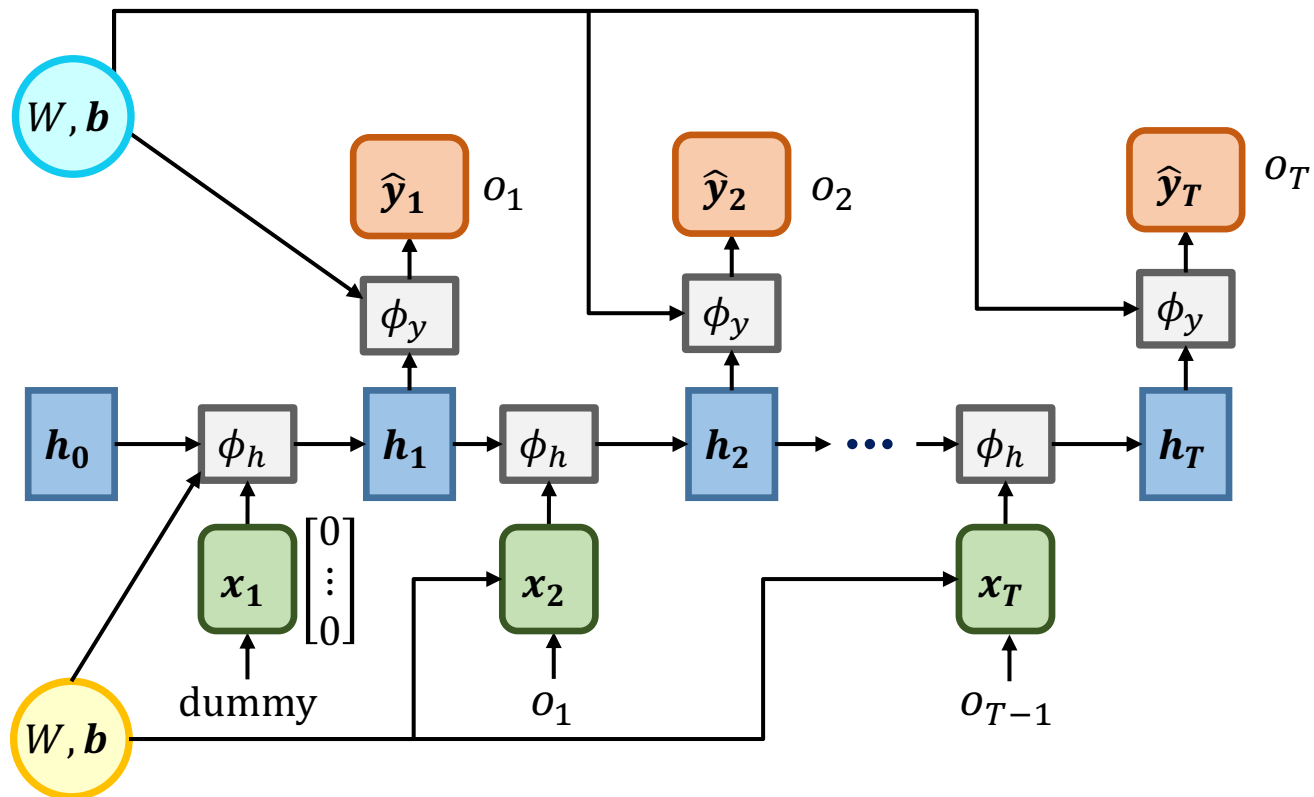
Computational Graph - revisited




Computational Graph - revisited



Computational Graph - revisited



Computational Graph - revisited



Training and Prediction using DKT

- Training: gradient descent
 - Prediction: compute inference in the network (see computational graph)
-

RNNs – Specifying Parameters

```
[ ] # Specify the model hyperparameters. Full descriptions included in the demo notebook!  
params = {}  
  
params['batch_size'] = 32  
params['mask_value'] = -1.0  
params['verbose'] = 1  
params['best_model_weights'] = 'weights/bestmodel'  
params['optimizer'] = 'adam'  
params['recurrent_units'] = 16  
params['epochs'] = 20  
params['dropout_rate'] = 0.1
```

RNNs – Tuning hyperparameters

```
[ ] # Specify the model hyperparameters. Full descriptions included in the demo notebook!  
params = {}  
  
params['batch_size'] = 32  
params['mask_value'] = -1.0  
params['verbose'] = 1  
params['best_model_weights'] = 'weights/bestmodel'  
params['optimizer'] = 'adam'  
params['recurrent_units'] = 16  
params['epochs'] = 20  
params['dropout_rate'] = 0.1
```


RNNs – Tuning hyperparameters

- Optimal number of epochs can be found using callbacks
 - Other parameters can be tuned using for example:
 - a) Train-Validation-Test split
 - b) Train-Test split, using a k-fold cross validation on the training data to determine the optimal parameters
-


Today – Recurrent Neural Networks

- Parameters and hyperparameter tuning
 - **Different architectures**
 - Different tasks:
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression
-

Recurrent Neural Network




Vanilla RNN - revisited




$$\mathbf{h}_t = \tanh(W_{hx}\mathbf{x}_t + W_{hh}\mathbf{h}_{t-1} + \mathbf{b}_h)$$

Long-Short Term Memory Network (LSTM)



- Two states:
 - Hidden state h_{t-1}
 - Cell state c_{t-1}

Long-Short Term Memory Network (LSTM)



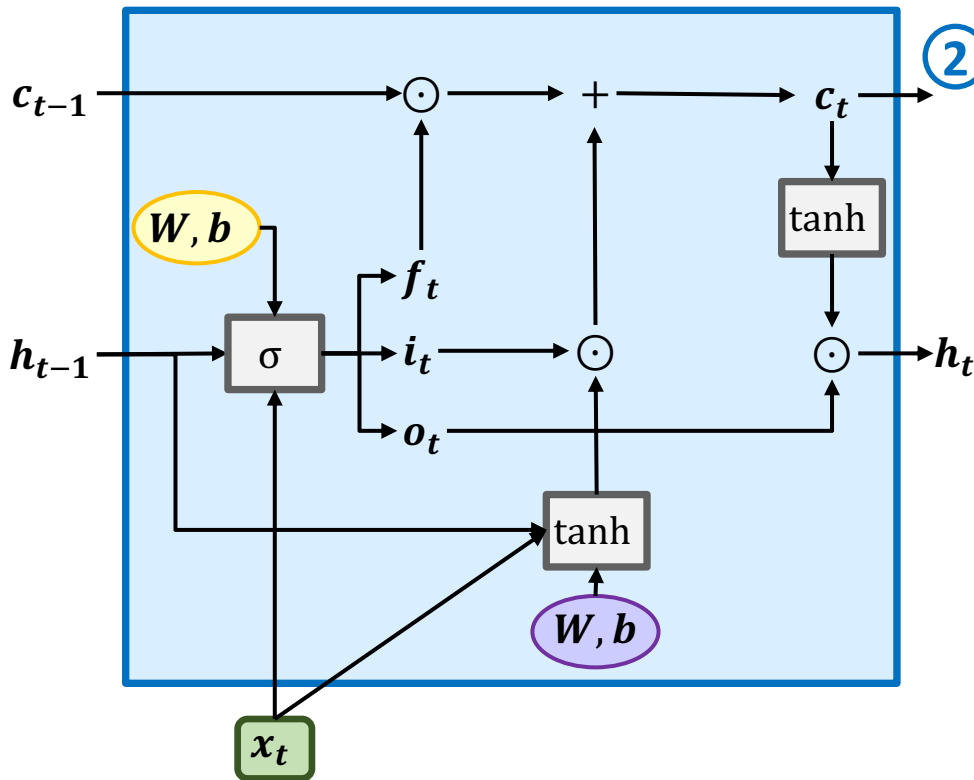
- ① Updating the gates:
- f forget gate: whether to erase cell
 - i input gate: whether to write to cell
 - o output gate: how much to reveal cell

$$f_t = \sigma(W_{fx}x_t + W_{fh}h_{t-1} + b_f)$$

$$i_t = \sigma(W_{ix}x_t + W_{ih}h_{t-1} + b_i)$$

$$o_t = \sigma(W_{ox}x_t + W_{oh}h_{t-1} + b_o)$$

Long-Short Term Memory Network (LSTM)




② Updating the states:

$$c_t = f_t \odot c_{t-1} + i_t \odot \tanh(W_{cx}x_t + W_{ch}h_{t-1} + b_c)$$


$$h_t = o_t \odot \tanh(c_t)$$

Gated Recurrent Units (GRU)



- Only one state (got rid of cell):
 - Hidden state h_{t-1}

Gated Recurrent Units (GRU)

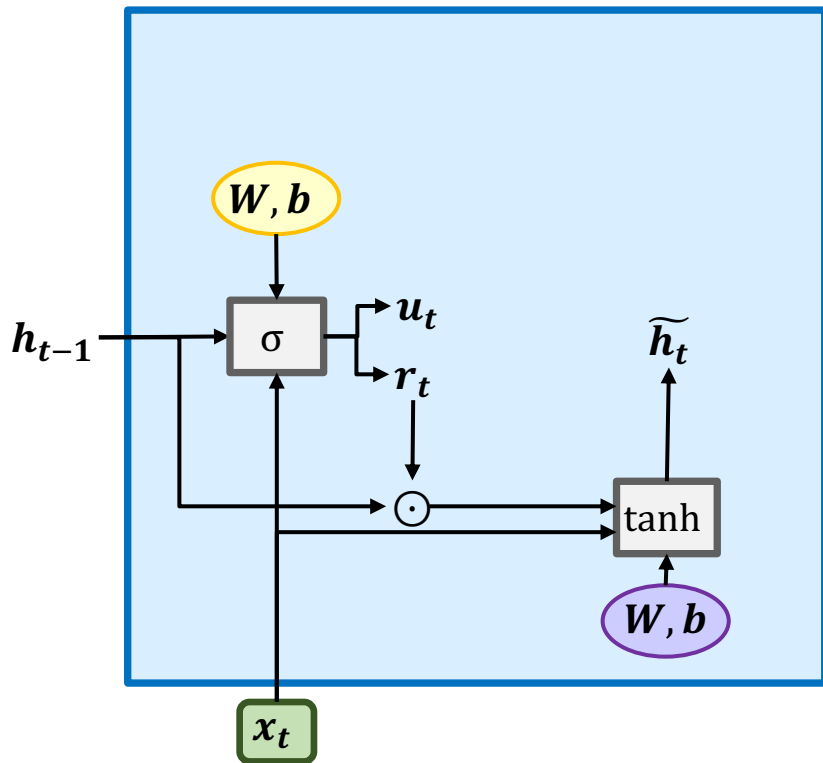


- ① Updating the gates:
- r reset gate: how much of the previous state to remember
 - u update gate: how much of the new state is just a copy of the old state

$$r_t = \sigma(W_{rx}x_t + W_{th}h_{t-1} + b_r)$$

$$u_t = \sigma(W_{ux}x_t + W_{uh}h_{t-1} + b_u)$$


Gated Recurrent Units (GRU)



② Get candidate hidden state:

$$\tilde{h}_t = \tanh(W_{hx}x_t + W_{ht}(r_t \odot h_{t-1}) + b_h)$$


Gated Recurrent Units (GRU)



③ Updating the state:

$$h_t = u_t \odot h_{t-1} + (1 - u_t) \odot \tilde{h}_t$$

Same input/output – different architectures



Today – Recurrent Neural Networks


- Parameters and hyperparameter tuning
 - Different architectures
 - **Different tasks:**
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression
-

Today – Recurrent Neural Networks

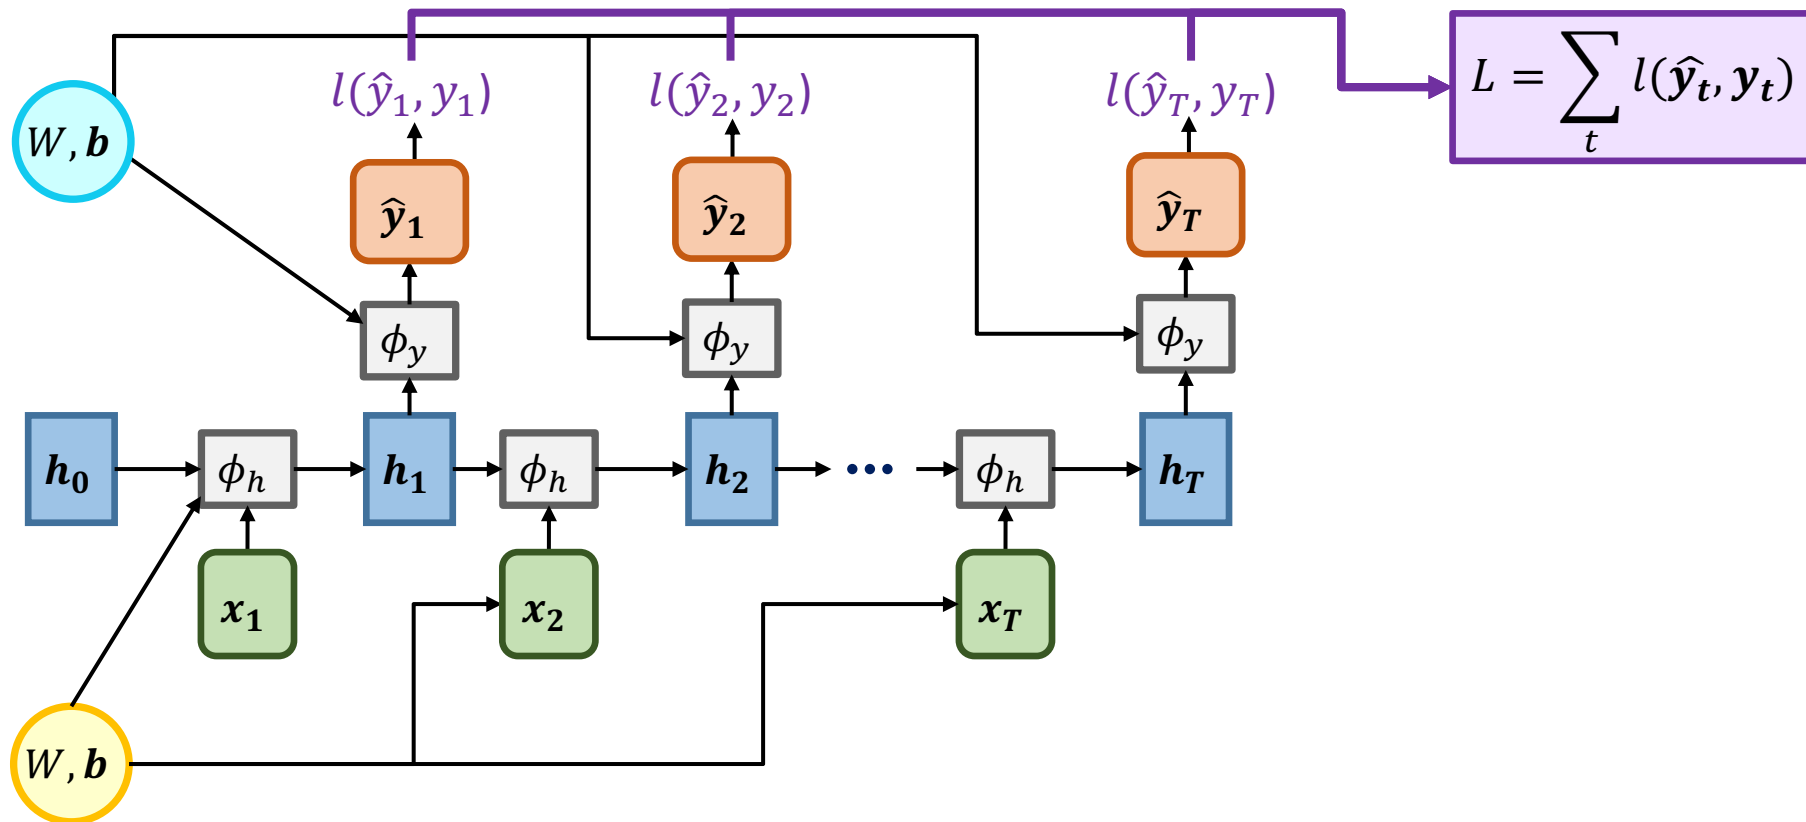
- Parameters and hyperparameter tuning
 - Different architectures
 - **Different tasks:**
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression
-

Many-to-many aka the Tracing Task

- Prediction of a target variable o_t at each time step t

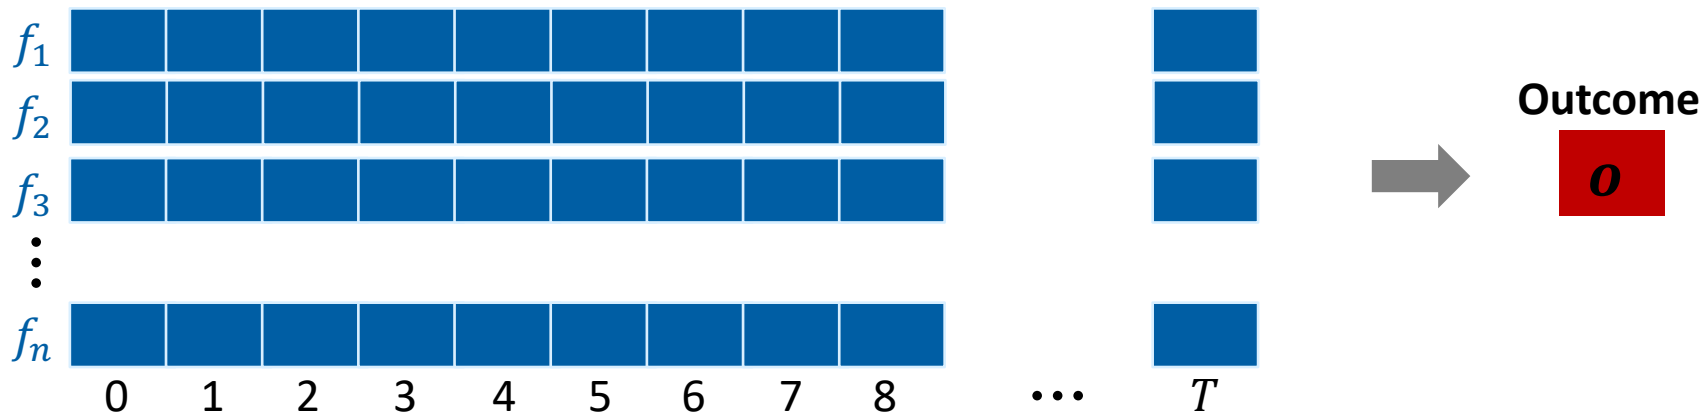


Computational Graph – Many-to-many

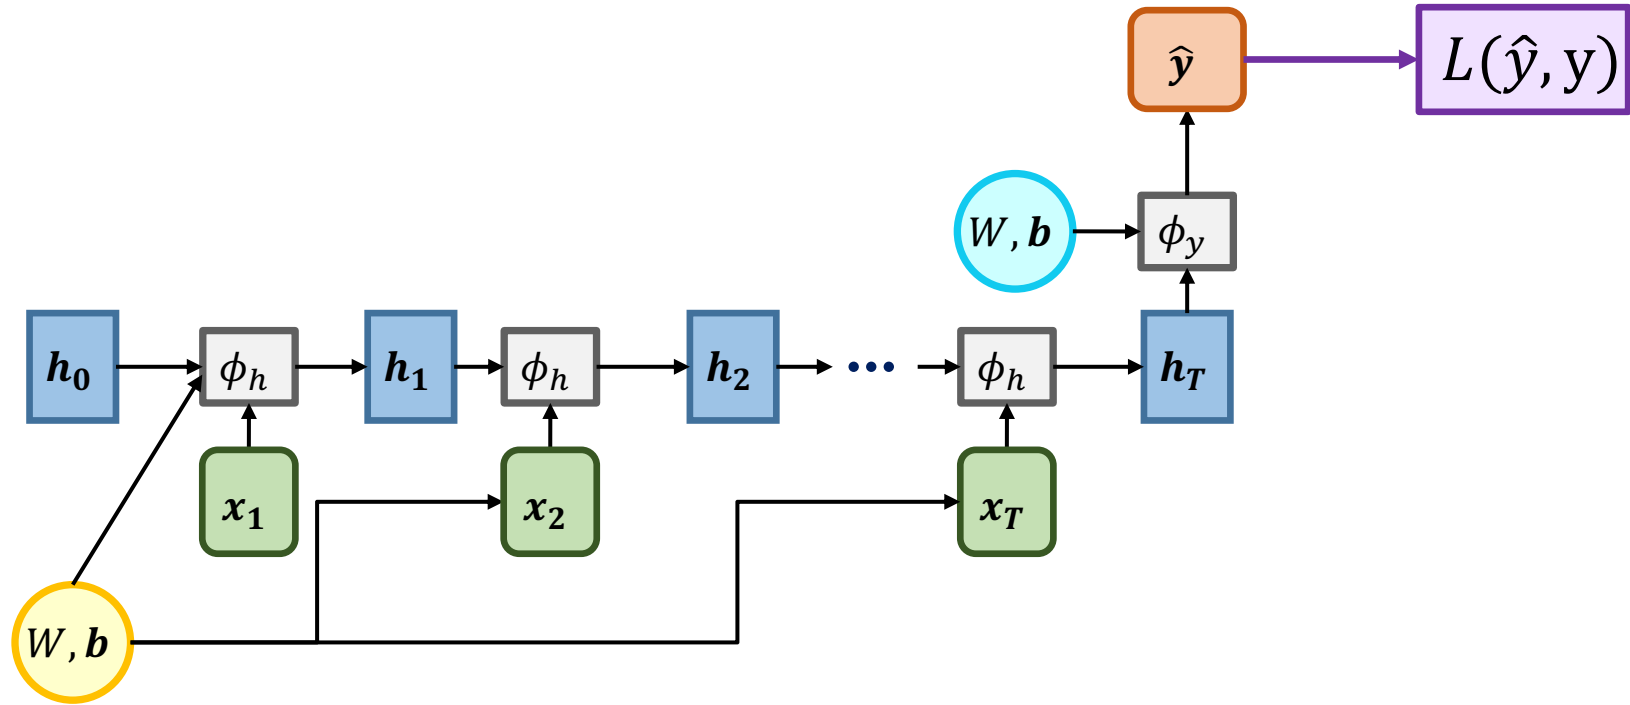


Many-to-one aka the Time-Series Prediction Task

- Prediction of a target variable o after $t \leq T$ time steps, where T is the total number of time steps



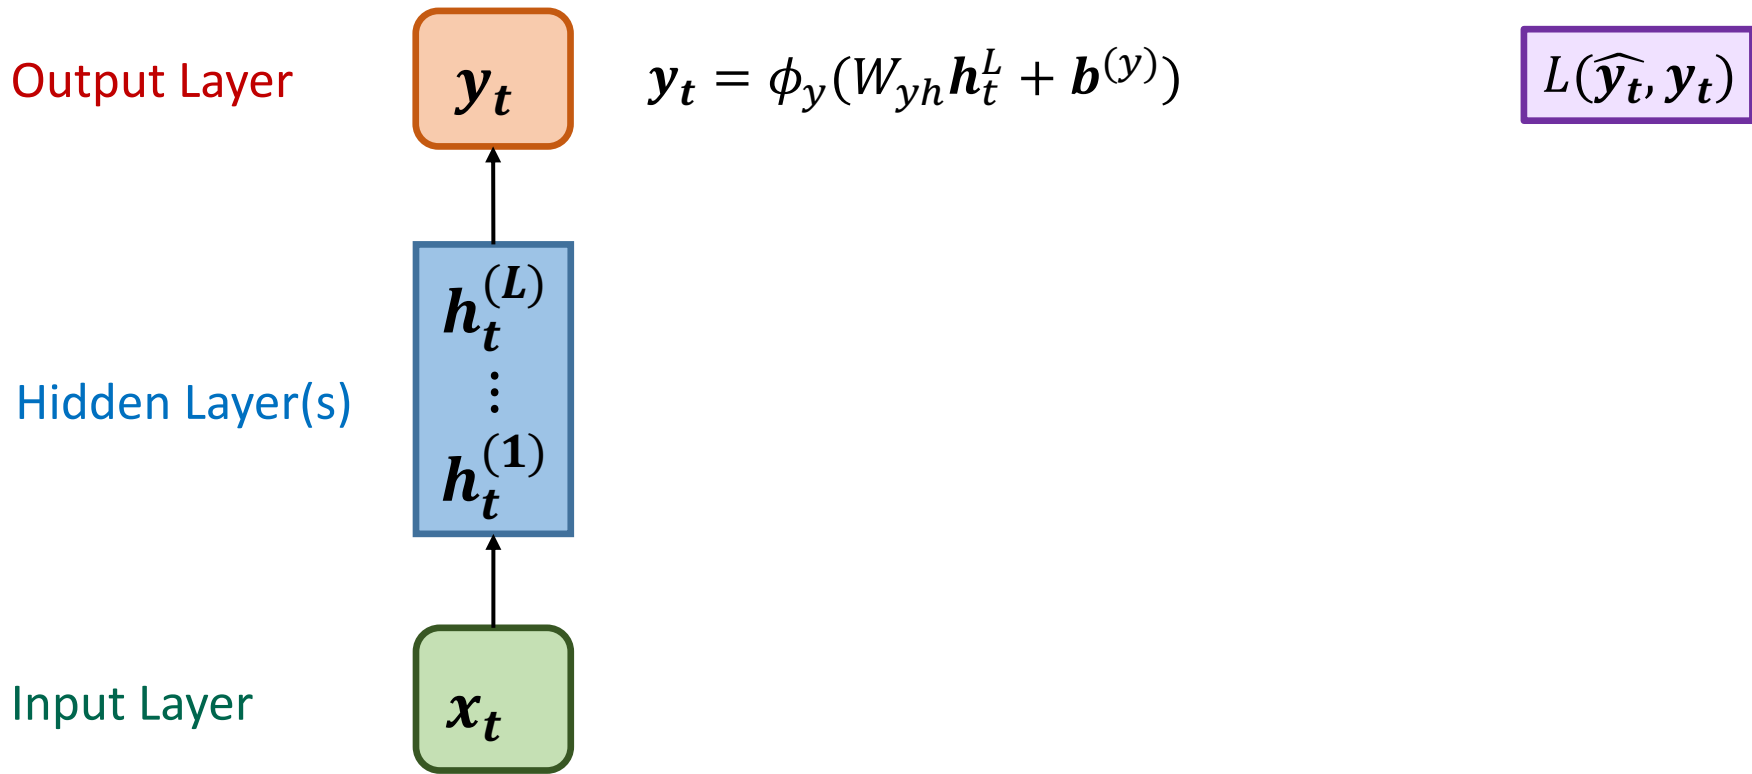
Computational Graph – Many-to-one



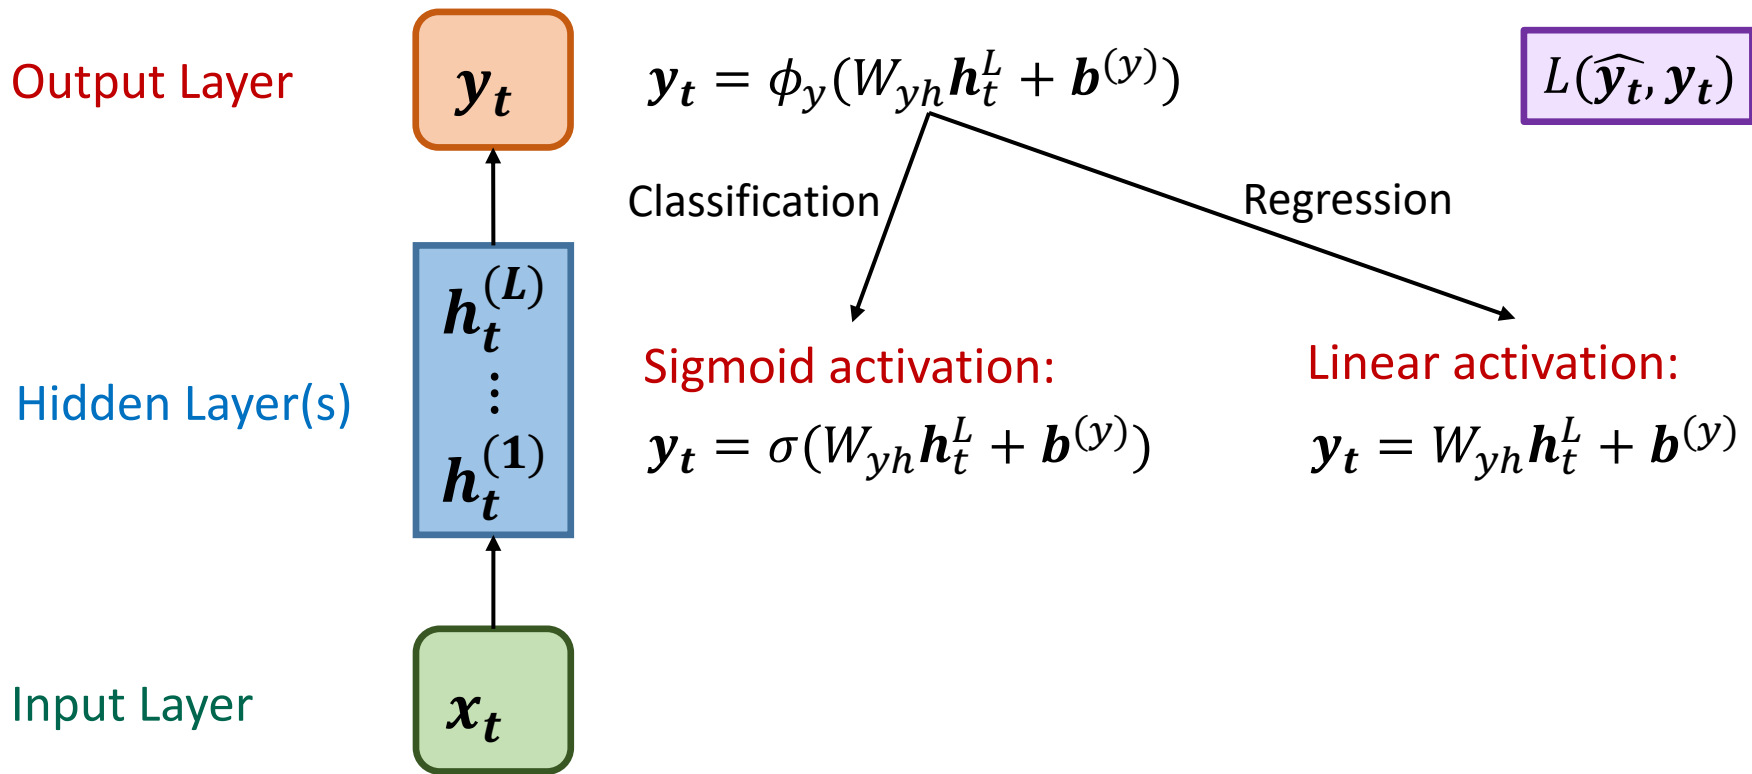
Today – Recurrent Neural Networks

- Parameters and hyperparameter tuning
 - Different architectures
 - **Different tasks:**
 - “Many-to-many” versus “Many-to-one”
 - **Classification versus Regression**
-

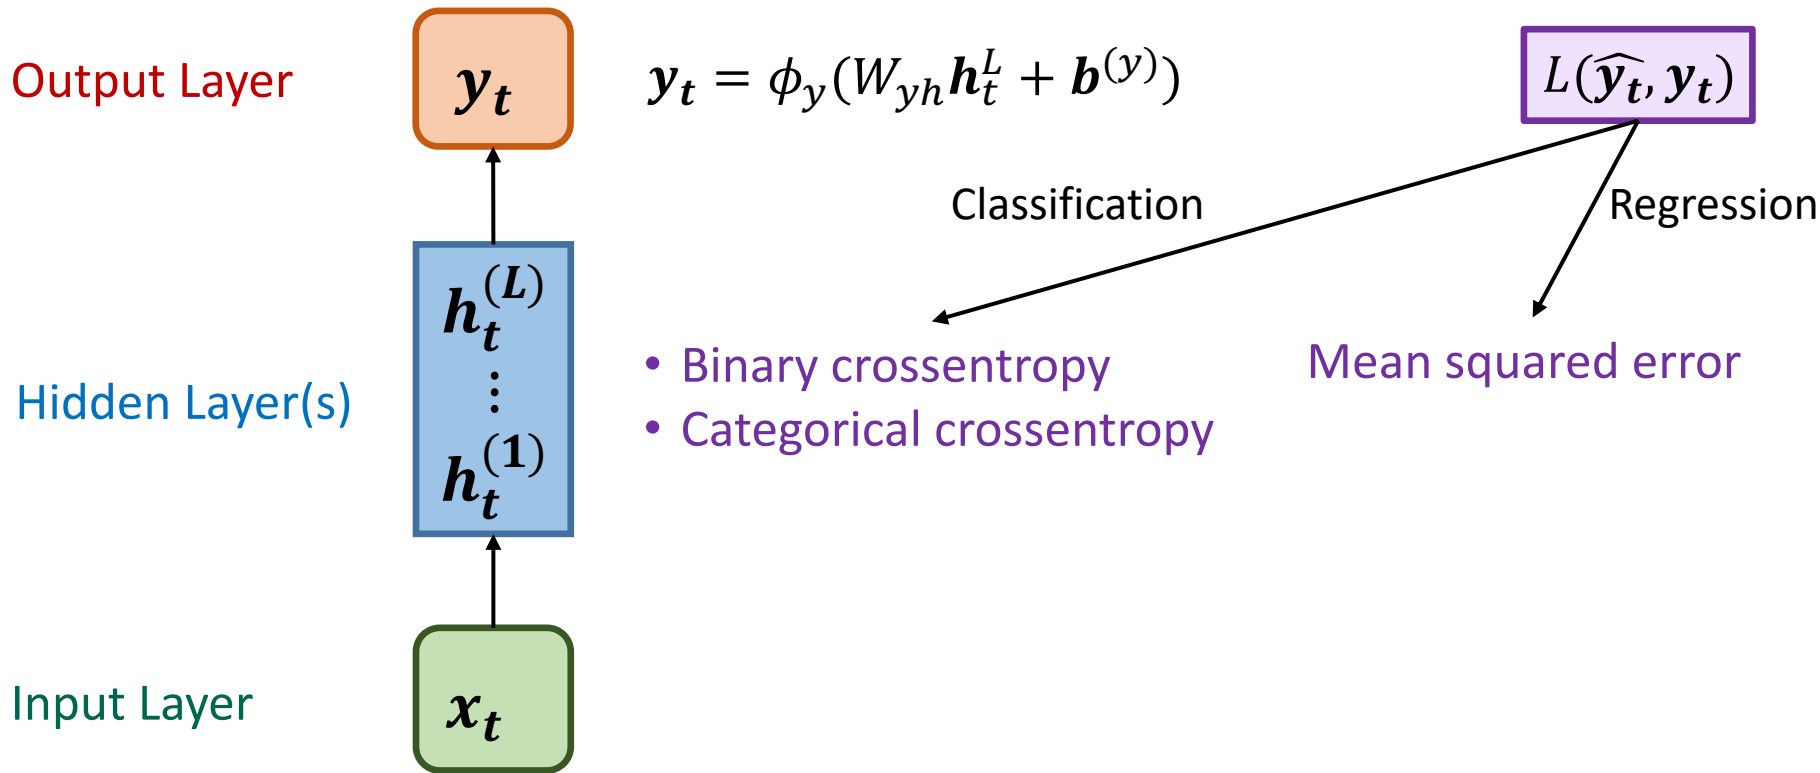
Classification vs. Regression



Classification vs. Regression: Output Layer



Classification vs. Regression: Training Loss



Your Turn

- Given:
 - Data from a MOOC
 - An LSTM for predicting quiz performance of a student for every week of the course (tracing task)
 - Your Task:
 - 1) Adjust the `create_model` function in order to predict pass/fail after 5 weeks of the course (time series prediction task) and send us the binary accuracy + AUC
 - Hint 1: `return_sequences=False`
 - Hint 2: what does `TimeDistributed(...)` do?
 - 2) Tune hyperparameters of your choice and send us binary accuracy and AUC
-

Summary

- Parameters and hyperparameter tuning
 - Different architectures
 - Different tasks:
 - “Many-to-many” versus “Many-to-one”
 - Classification versus Regression
-

Lab Session on April 13

- First part (8.15-9.00): hands-on tutorial on RNNs
 - Second part (starting at 9.15): office hours for questions regarding the project
-